🧪 Candidate Challenge: Mini Evidence Retrieval Task

Background

We are exploring AI systems similar to OpenEvidence, which analyze scientific and business data to provide **structured, source-backed insights**.

Your task is to demonstrate how you would **ingest, search, and summarize evidence from documents**.

The Challenge (2–3 hours)

1. Dataset

We provide you with **10–15 scientific abstracts** (or you may choose a small set of PubMed abstracts, news articles, or PDFs of your choice).

2. Task

Build a **small pipeline** that:

1. **Ingests the documents** (PDF, plain text, or abstracts).
2. **Indexes them** for semantic search (e.g., embeddings + vector database, or equivalent).
3. **Allows a user query** (e.g., *"What are common side effects of diabetes medication?"*).
4. **Returns**:
    - A short AI-generated answer (2–3 sentences).
    - A list of 2–3 source documents with citations.

3. Deliverables

- **Code** (Python, Node.js, or another language you prefer).
- **Brief README** explaining:
    - How your pipeline works.
    - How to run it.
    - What design choices you made (libraries, architecture, assumptions).

(You do **not** need a polished UI)

Evaluation Criteria

We will look at:

- ✅ **Data Handling** – Can you parse and index the documents effectively?
- ✅ **Retrieval Quality** – Does your system return relevant sources?
- ✅ **AI Integration** – Is the generated summary grounded in the retrieved documents?
- ✅ **Clarity** – Is your code and README understandable?
- ✅ **Practical Thinking** – Did you make reasonable trade-offs given the short time?

👉 This challenge is intentionally scoped small — we don't expect a full OpenEvidence clone. The goal is to see **how you think about evidence-based AI pipelines**.